

Prime Throughput Performance Evaluation

Report 1

Source Code and GitHub Repositories

The files used for this evaluation can be found at the following repository:

<https://github.com/sam444555/cs2910>

The original Prime code used is available at the following Johns Hopkins DSN repository:

<https://github.com/jhu-dsn/prime>

Host System Specifications

The throughput evaluation was performed locally using Docker on a system with the following specifications:

- OS: Windows 11 Education, Version 25H2 (Build 26200.9457) [as of 9/24/2026]
- Intel Core i7-14700KF
 - 20 cores: 8 P-cores, 12 E-cores
 - 3.4 GHz base clock speed, up to ~5.6 GHz boost
- Memory: 32 GB DDR5 @ 4800 MT/s

WINDOWS WARNING:

- Microsoft's Hyper-V virtualization technology enables Linux-based Docker containers to run on Windows via WSL2 (Windows Subsystem for Linux 2).
- WSL2 provides a Linux environment (Ubuntu 22.04 LTS) within a virtual machine hosted on Windows where Docker containers share a common Linux kernel.
- Following a Windows update, the throughput numbers recorded here dropped significantly
 - Ex: Unrestricted Network Test, 25 emulated clients, 10k total updates
 - Pre-update: ~1042 updates/sec
 - Post-update: ~42 updates/sec
- After troubleshooting, the cause turned out to be Hyper-V's TimeSync incorrectly applying corrections to WSL's wall clock, pushing it 5.4-5.8 seconds ahead.

- A fix for this issue is available on the GitHub. I would highly recommend first running this deployment on a native Linux OS to avoid the headaches that come with Windows and its virtualization.

Docker Configuration

The Docker files & setup guide can be found at:

<https://github.com/sam444555/cs2910/tree/main/prime/docker>

6 Prime replicas were deployed across six Docker containers, satisfying Prime's minimum replica requirement of $3f + 2k + 1$ ($f=1, k=1$) as defined in

<https://github.com/sam444555/cs2910/blob/main/prime/src/def.h> by the **NUM_SERVERS** macro.

Each container ran a Prime replica and a Spines process. One container additionally ran the driver program used to emulate additional clients. Each replica was assigned a unique positive integer identifier starting at 1.

For the throughput testing, the containers were named **prime1, prime2, prime3, prime4, prime5, prime6**, with the driver program running on **prime1**. The six containers communicated over a custom Docker bridge network. The complete Docker topology can be found in Figure 1. The containers use **Ubuntu 22.04** as their operating system.

Docker Topology (6 Replicas)

Prime + Spines with Driver (prime1)




Figure 1. Docker network topology

Driver Setup

The driver program used for the testing can be located at:

<https://github.com/sam444555/cs2910/blob/main/prime/src/driver.c>

The original **driver.c** program did not include throughput measurement, so it had to be modified to include this functionality. Throughput is calculated using the following

$$\text{equation: } \textit{Throughput} = \frac{\textit{Total Updates Completed}}{\textit{Elapsed Time (seconds)}}$$

Elapsed time is calculated using the throughput stopwatch (throughput_sw in driver.c) which starts in Run_Client() and stops in Process_Message() once all updates have been successfully processed.

Total Updates Completed is represented by the needed_count variable and is input into the program using the -c (benchmark count) flag.

Another modification made to driver.c was to pass the number of emulated clients using the -n flag, replacing the macro definition in order to eliminate the need to recompile the

program every time this value changes. It is now stored in the num_clients_to_emulate variable.

Additional changes were made; however, these are discussed later in the document to provide the appropriate context.

Running the driver program:

`./driver -l <LOCAL_IP> -i <CLIENT_ID> -s <SERVER_ID> -c <UPDATE_COUNT> -n <EMULATED_CLIENTS>`

- LOCAL_IP: IP address of the machine the driver program is running on
- CLIENT_ID: A positive integer starting from 1 which uniquely identifies the client process
- SERVER_ID: The ID of the prime replica the client process is communicating with.
- UPDATE_COUNT: The total number of updates the driver program will send
- EMULATED_CLIENTS: Number of clients being emulated by the driver program, which also determines the maximum number of outstanding updates




Figure 2. driver.c control flow chart

Testing: No Throughput/Latency Constraints

In the original driver's code, once the number of emulated clients exceeded ~20, the driver would unpredictably stall indefinitely across runs. In the driver's `Send_Update()` function, its `sendto()` call was blocking. At approximately the ~20-client threshold, the `sendto()` call would sporadically block and fail to return control to the driver, causing the process to stall permanently.

I changed the `sendto()` call to nonblocking and created an error handler for failed sends. Upon detecting a failure to send, the handler would roll back changes made by the `Send_Update()` function (specifically the `time_stamp` increment which is used for driver update sequence numbers), record the failure to send (tracked by the global variable `failed_sends`), and return control to the event handler (`E_handle_events()`).

This modification reduced the number of stalls, but did not eliminate them entirely. Without a deeper understanding of Prime, I cannot say with 100% certainty what the cause of this behavior is. However, based on debugging and testing, my initial leading theory was that the stalls were due to the driver's socket's outgoing queue becoming full due to higher throughput and Prime not draining it fast enough.


When `Send_Update()` performs a successful send (i.e. `sendto()` does not return an error), the function increments the number of outstanding updates (the number of updates sent to Prime for which a response has not yet been received). During stalls, however, the value of outstanding updates was nonzero, meaning the driver "thought" it had successfully sent >0 updates. Under normal operation, the event handler would eventually call `Process_Message()` when Prime responds to these outstanding messages. But this behavior never occurred, resulting in the freeze.

Therefore, my initial theory is likely incorrect, and the most likely explanation for the stalling is Prime's queuing behavior and its decision to reject packets when it becomes saturated. The driver's current socket configuration for IPC uses a Unix-domain datagram socket, which does not guarantee data delivery. There is an option to disable IPC communication in `prime/src/def.h` by setting the `USE_IPC_CLIENT` macro to 0. However, I encountered difficulties establishing the TCP connection between the driver and Prime, so this configuration was not evaluated further.


The following tests were conducted with no throughput constraint applied (e.g. no latency or bandwidth limitations). These results serve as the best-case baseline for Prime's throughput performance.

No Throughput/Latency Constraint Tests (total updates = 10k)			
Number of Emulated Clients	Time Elapsed (secs)	Total Failed Sends	Throughput (updates/sec)
1	201.48	0	49.63
5	40.3	0	248.12
10	20.21	0	494.69
15	13.55	70	737.93
20	10.09	240	990.89
25	9.6	1121	1041.58
30	9.62	3561	1039.97
50	9.63	9968	1038.55
100	9.71	9961	1030.04

Throughput vs. Number of Emulated Clients
 (Total Updates = 10k, No Throughput/Latency Constraints)



Total Failed Sends vs. Number of Emulated Clients
(Total Updates = 10k, No Throughput/Latency Constraints)



The throughput peaked in the range of 1,000-1,040 updates/sec. It began to stagnate once the number of emulated clients exceeded ~25. Anything beyond this value resulted in no additional benefit and led to an explosion in cumulative sending failures.

Therefore, I cannot conclusively say that Prime's "best-case" baseline is ~1,040 updates/sec due to the instability observed as the system approaches this level of throughput. If my theory that the instability stems from Prime itself discarding updates due to overflow is correct, employing a transport protocol with reliable delivery guarantees (such as TCP) could potentially resolve this issue. However, this additional overhead would likely result in a lower throughput. The significance of the reduction is difficult to determine given I was unable to get TCP to function.

In practice, the highest consistently stable throughput observed during testing occurred with 15-16 emulated clients, corresponding to roughly 740-800 updates/second.

Testing: Bandwidth & Latency Constraints

0ms (~.067 ms avg) RTT, 100 Mbps Link Speed (total updates = 10k)			
Number of Emulated Clients	Time Elapsed (secs)	Total Failed Sends	Throughput (updates/sec)
1	207.32	0	48.24
5	41.38	0	241.64
10	21.93	0	456
15	17.03	81	587.1
20	13.86	173	721.6
25	14.03	684	712.57
30	14.04	2684	712.32
50	13.26	9968	753.96
100	13.42	9968	744.93

1 ms RTT, 100 Mbps Link Speed (total updates = 10k)			
Number of Emulated Clients	Time Elapsed (secs)	Total Failed Sends	Throughput (updates/sec)
1	215.32	0	46.44
5	47.22	0	211.76
10	29.3	0	341.28
15	23.07	61	433.44
20	19.94	167	501.5
25	17.48	932	572.14
30	18.8	3113	531.86
50	19.05	9968	524.89
100	19.2	9973	520.97

5 ms RTT, 100 Mbps Link Speed (total updates = 10k)			
Number of Emulated Clients	Time Elapsed (secs)	Total Failed Sends	Throughput (updates/sec)
1	401.3	0	24.29
5	80.25	0	124.61
10	40.1	0	249.39
15	26.81	4	372.99
20	20.14	12	496.45
25	18.36	184	544.59
30	18.34	743	545.24
50	18.44	9975	542.17
100	18.36	9963	544.72


20 ms RTT, 100 Mbps Link Speed (total updates = 10k)			
Number of Emulated Clients	Time Elapsed (secs)	Total Failed Sends	Throughput (updates/sec)
1	840.5	0	11.9
5	170.27	0	58.73
10	83.81	0	119.31
15	54.6	0	183.13
20	42.52	0	235.18
25	35.44	3	282.14
30	35.68	48	280.29
50	35.4	9959	281.98
100	35.69	9963	280.2

40 ms RTT, 100 Mbps Link Speed (total updates = 10k)			
Number of Emulated Clients	Time Elapsed (secs)	Total Failed Sends	Throughput (updates/sec)
1	Not Measured	Not Measured	Not Measured
5	293.41	0	34.08
10	147.07	0	67.99
15	97.5	1	102.56
20	74.21	0	134.75
25	61.93	6	161.48
30	61.7	98	162.08
50	59.8	9961	167.23
100	60.83	9965	164.38

The stalling issues occurred, albeit far less frequently, in these tests as well. A reliable delivery guarantee mechanism should be incorporated to ensure consistent system stability. The resulting latency of this mechanism should also be considered given its impact on throughput performance observed throughout testing.

Peak Throughput vs. Network RTT

(Total Updates = 10k, 100 Mbps Link Speed)



All Tests Graphed

